

Scalability and Map Reduce Programming

CS 144 Web Application

TA: Jin Wang

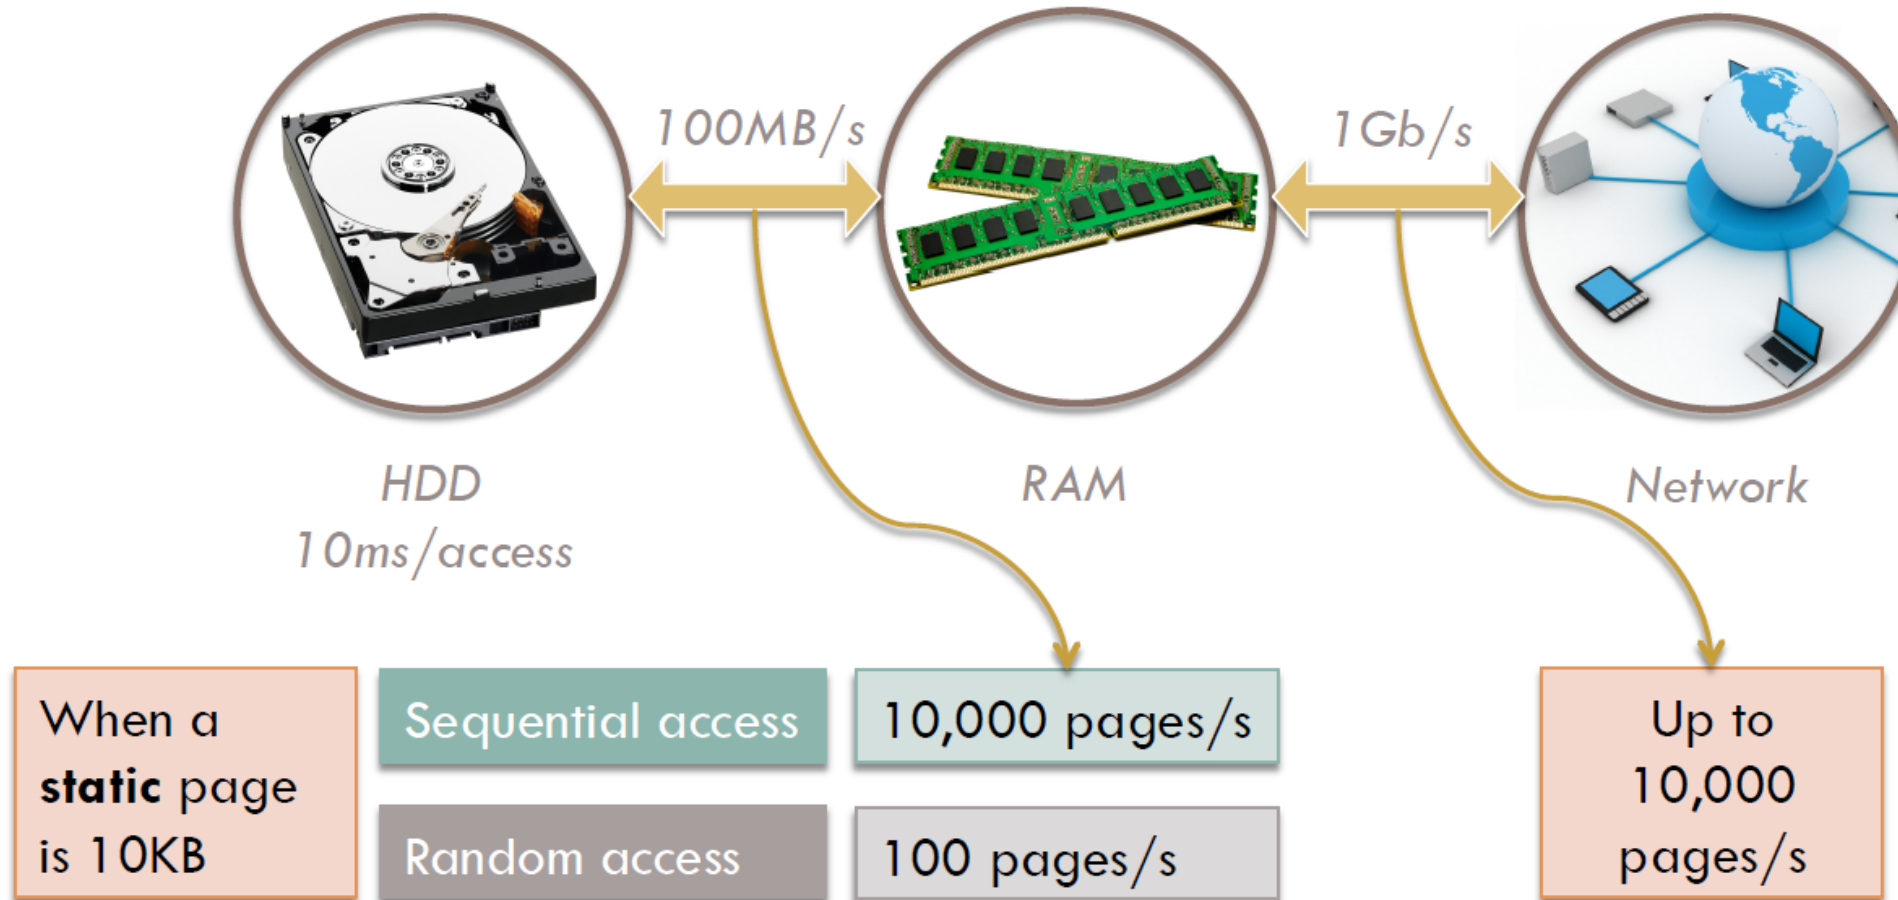
03/02/2018

Overview

- Review of foundations about scalability
- The Hadoop platform
- **Map Reduce Programming**

Scalability

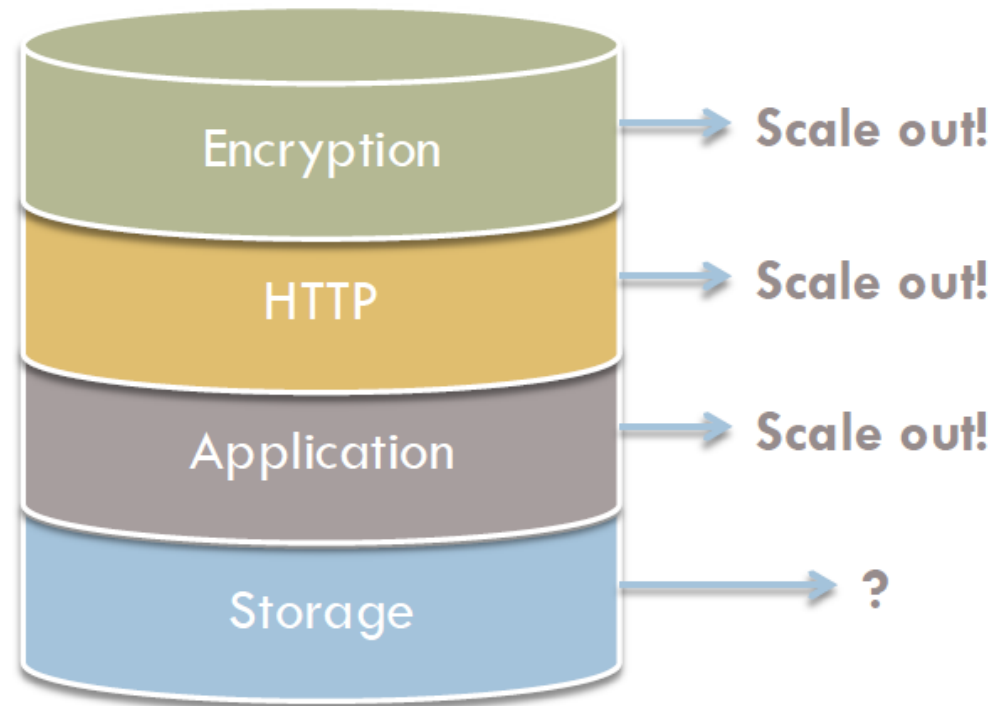
How to estimate our server capabilities:



Scaling Web Applications


- Scale out
- Scale up

Web Server Architecture



Scaling Web Applications


- The Load Balancer
 - TCP-NAT Request Distributor:
DNS Round Robin, or software



- In our projects, Apache Tomcat is both in the application and HTTP layers.

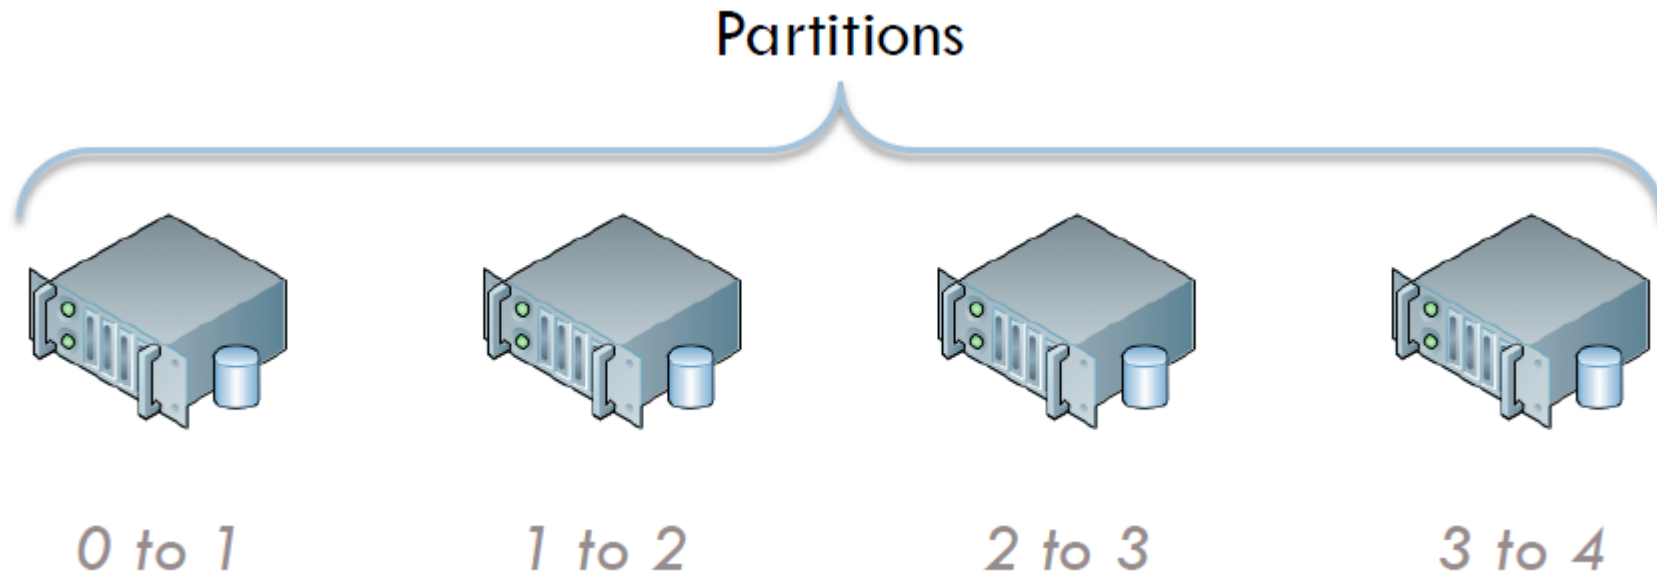
Scaling the Storage Layer

- Scenario 1: Read Only
 - Information doesn't change. Clients only read data.
 - Use **replication**





Scaling the Storage Layer

- *Scenario 2: Local Read/Write*
 - Reads and writes are scoped to individual users
 - Use **shard** or **partitioning**



Scaling the Storage Layer

- *Scenario 3: Global Read/Write*
 - Reads and writes are global, and all users can see everyone updates
 - Use **partitioning** and **replication**




Introduction to Hadoop

- Download from hadoop.apache.org
- To install locally, unzip and set JAVA_HOME
- Docs: hadoop.apache.org/common/docs/current
- Three ways to write jobs:
 - Java API
 - Hadoop Streaming (for Python, Perl, etc)
 - Pipes API (C++)

Note: The following slides are borrowed from Prof. Tyson Condie

Typical Hadoop Cluster



- 40 nodes/rack, 1000-4000 nodes in cluster
- 1 Gbps bandwidth in rack, 8 Gbps out of rack
- Node specs (Facebook):
8-16 cores, 32 GB RAM, 8×1.5 TB disks, no RAID

Typical Hadoop Cluster




Hadoop Components

- Distributed file system (HDFS)
 - Single namespace for entire cluster
 - Replicates data 3x for fault-tolerance
- MapReduce framework
 - Runs jobs submitted by users
 - Manages work distribution & fault-tolerance
 - Colocated with file system



Hadoop Distributed File System (HDFS)

- Files split into 128MB blocks
- Blocks replicated across several datanodes (often 3)
- Namenode stores metadata (file names, locations, etc)
- Optimized for large files, sequential reads
- Files are append-only



What is MapReduce?

- Programming model for data-intensive computing on commodity clusters
- Pioneered by Google
 - Processes 20 PB of data per day
- Popularized by Apache Hadoop project
 - Used by Facebook, Amazon, ...

What is MapReduce Used For?

- At Google:
 - Originally: Index building for Google Search
 - Article clustering for Google News
 - Statistical machine translation
- At Facebook:
 - Data mining
 - Ad optimization
 - Spam detection

MapReduce Programming Model

- Data type: key-value *records*

- Map function:

$$(K_{in} , V_{in}) \Rightarrow \text{list}(K_{inter} , V_{inter})$$

- Reduce function:

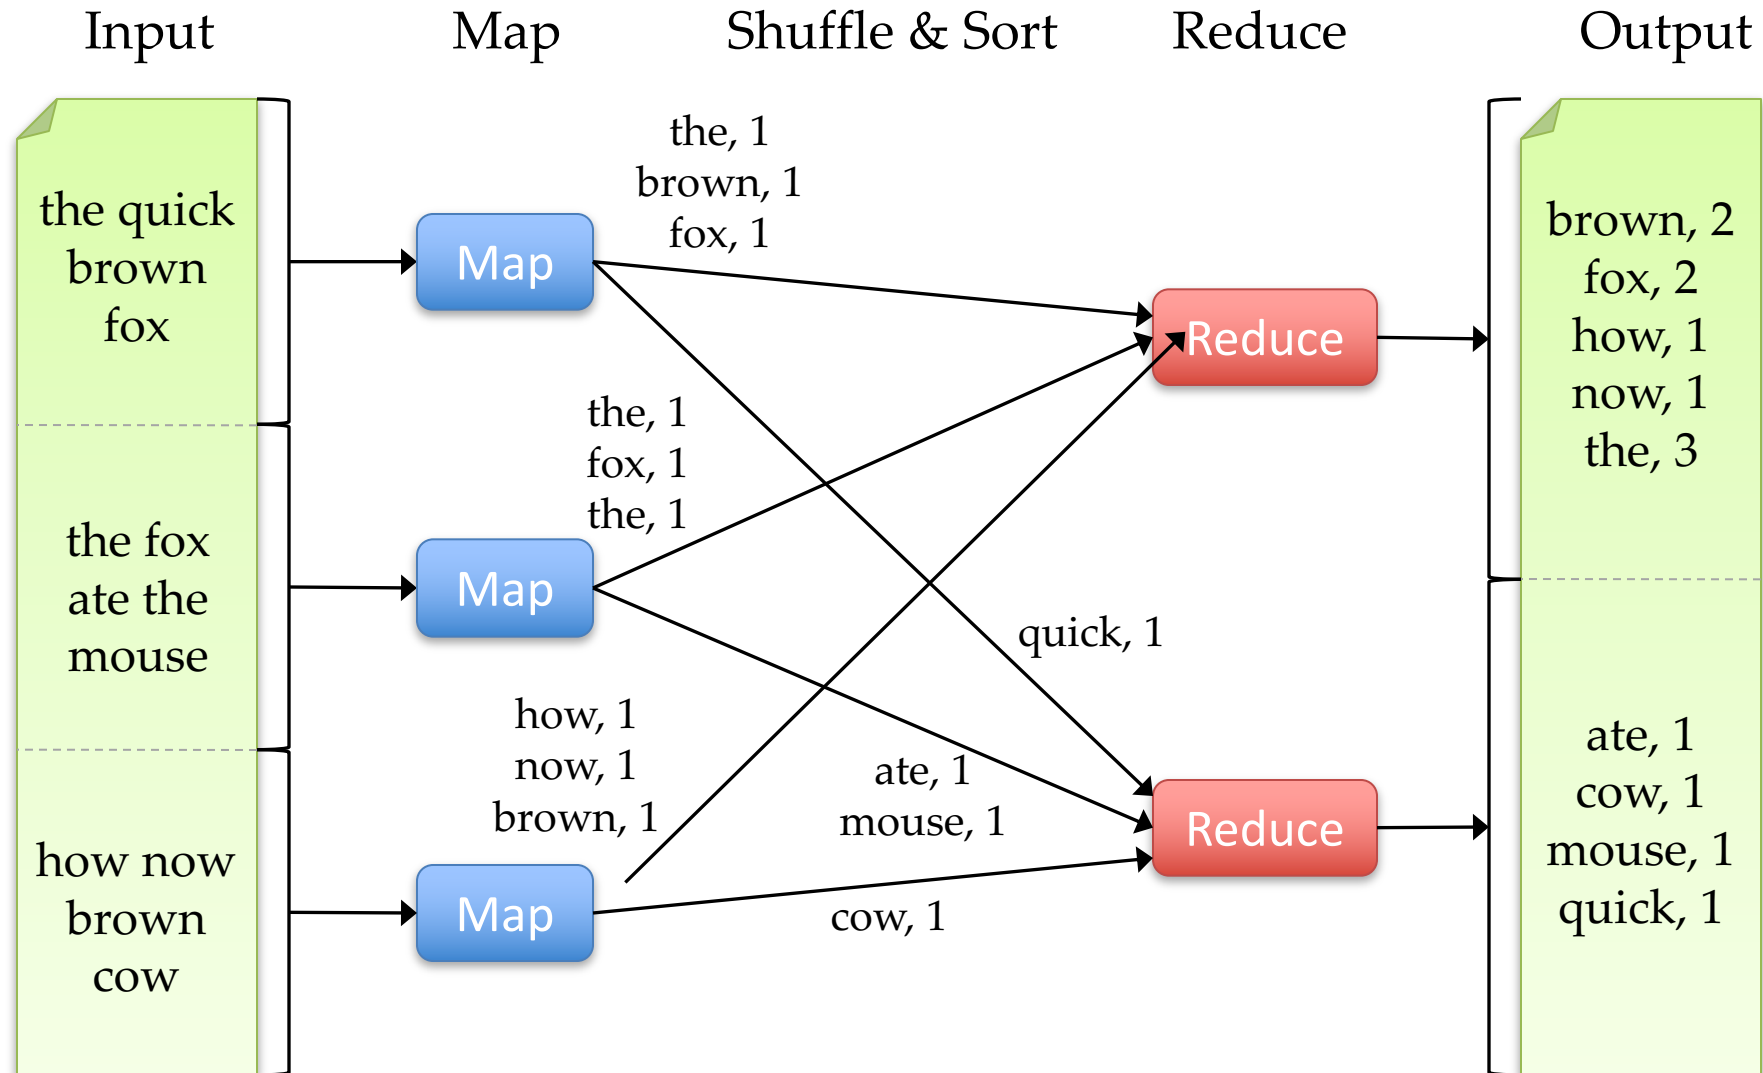
$$(K_{inter} , \text{list}(V_{inter})) \Rightarrow \text{list}(K_{out} , V_{out})$$

Example: Word Count

```
def mapper(line):  
    foreach word in line.split():  
        output(word, 1)
```

```
def reducer(key, values):  
    output(key, sum(values))
```

Word Count Execution

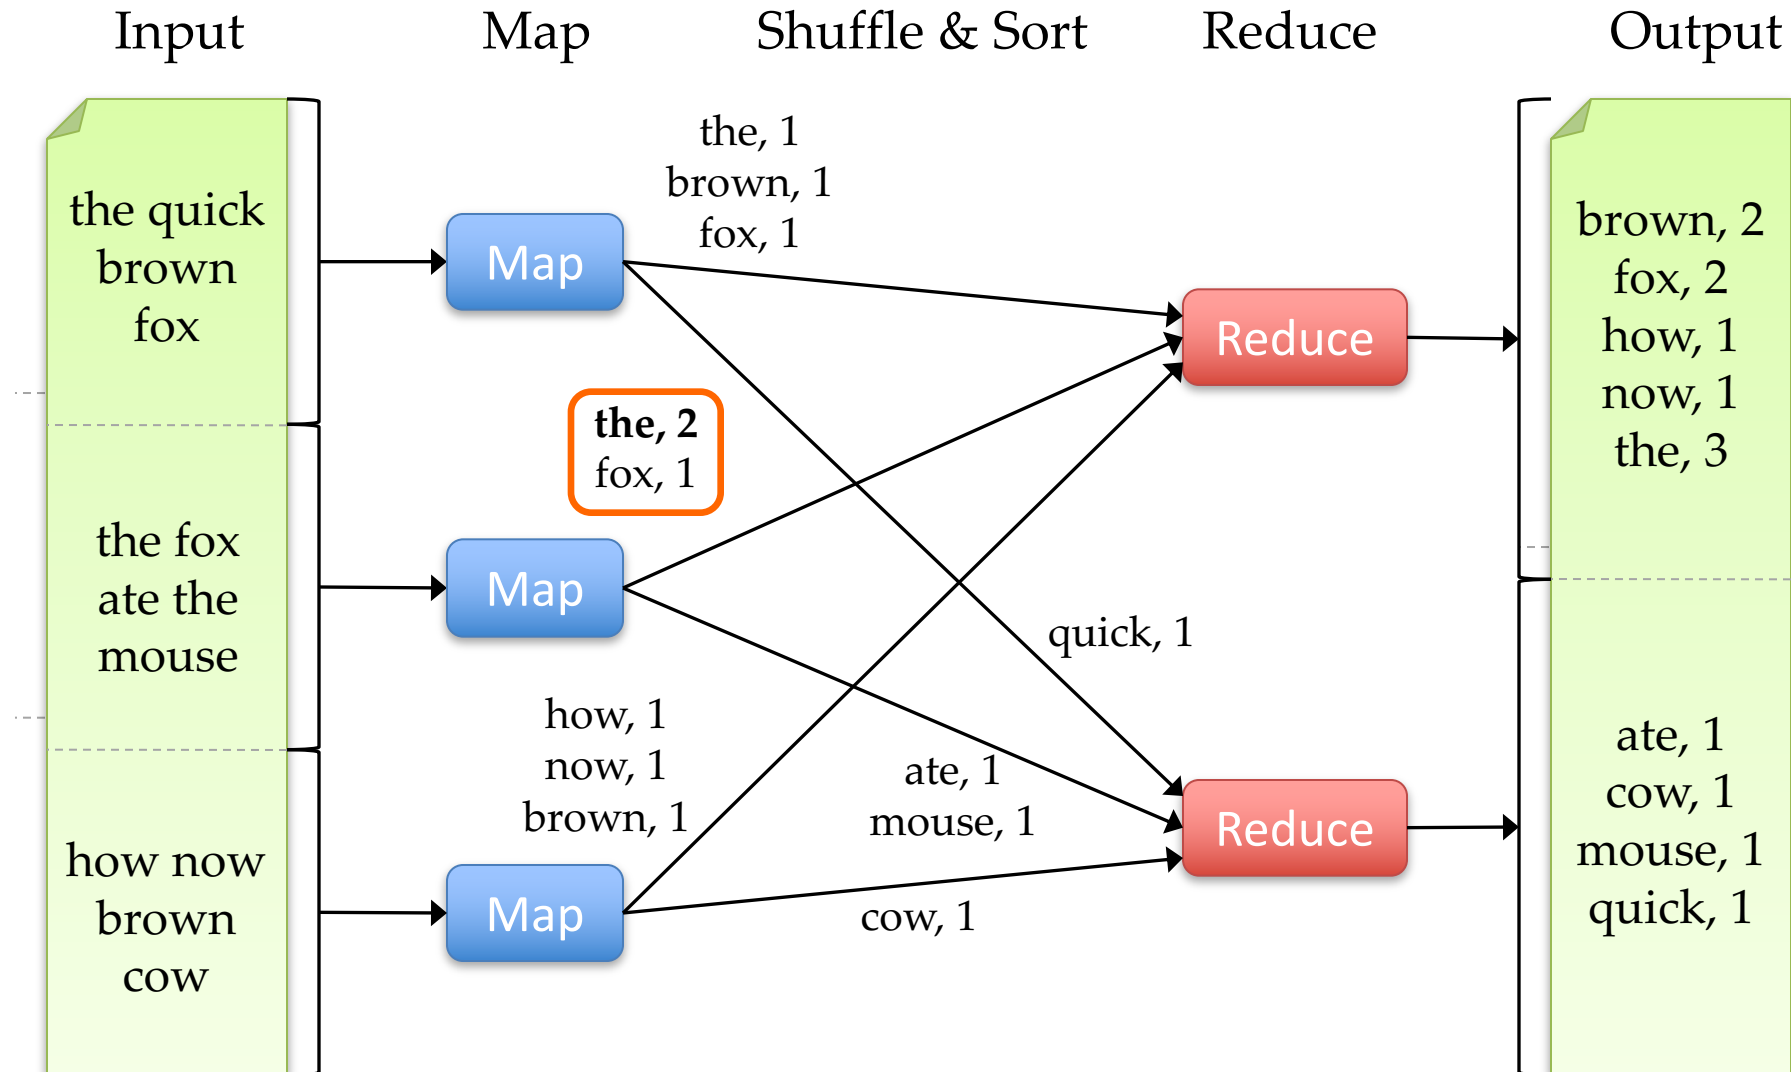


An Optimization: The Combiner

- Local reduce function for repeated keys produced by same map
- For associative ops. like sum, count, max
- Decreases amount of intermediate data
- Example: local counting for Word Count:

```
def combiner(key, values):  
    output(key, sum(values))
```

Word Count with Combiner



MapReduce Execution Details

- Mappers preferentially scheduled on same node or same rack as their input block
 - Minimize network use to improve performance
- Mappers save outputs to local disk before serving to reducers
 - Allows recovery if a reducer crashes
 - Allows running more reducers than # of nodes

Examples of Map Reduce Programming

1. Search

- **Input:** (lineNumber, line) records, a given pattern
- **Output:** lines matching the pattern

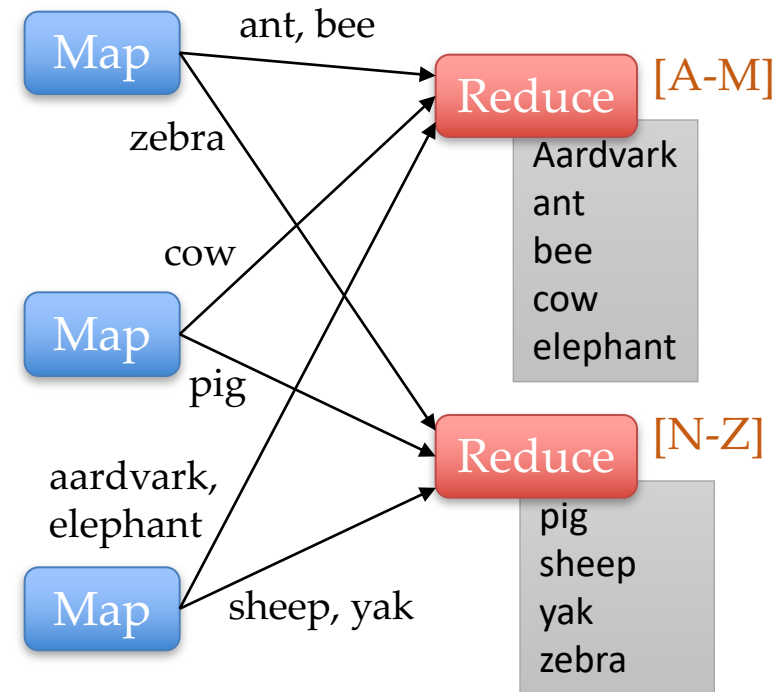
- **Map:**

```
if(line matches pattern):  
    output(line)
```

- **Reduce:** identity function
 - Alternative: no reducer (map-only job)

2. Sort

- **Input:** (key, value) records
- **Output:** same records, sorted by key
- **Map:** identity function
- **Reduce:** identify function
- **Trick:** Pick partitioning function p such that $k_1 < k_2 \Rightarrow p(k_1) < p(k_2)$



3. Inverted Index

- **Input:** (filename, text) records
- **Output:** list of files containing each word

- **Map:**

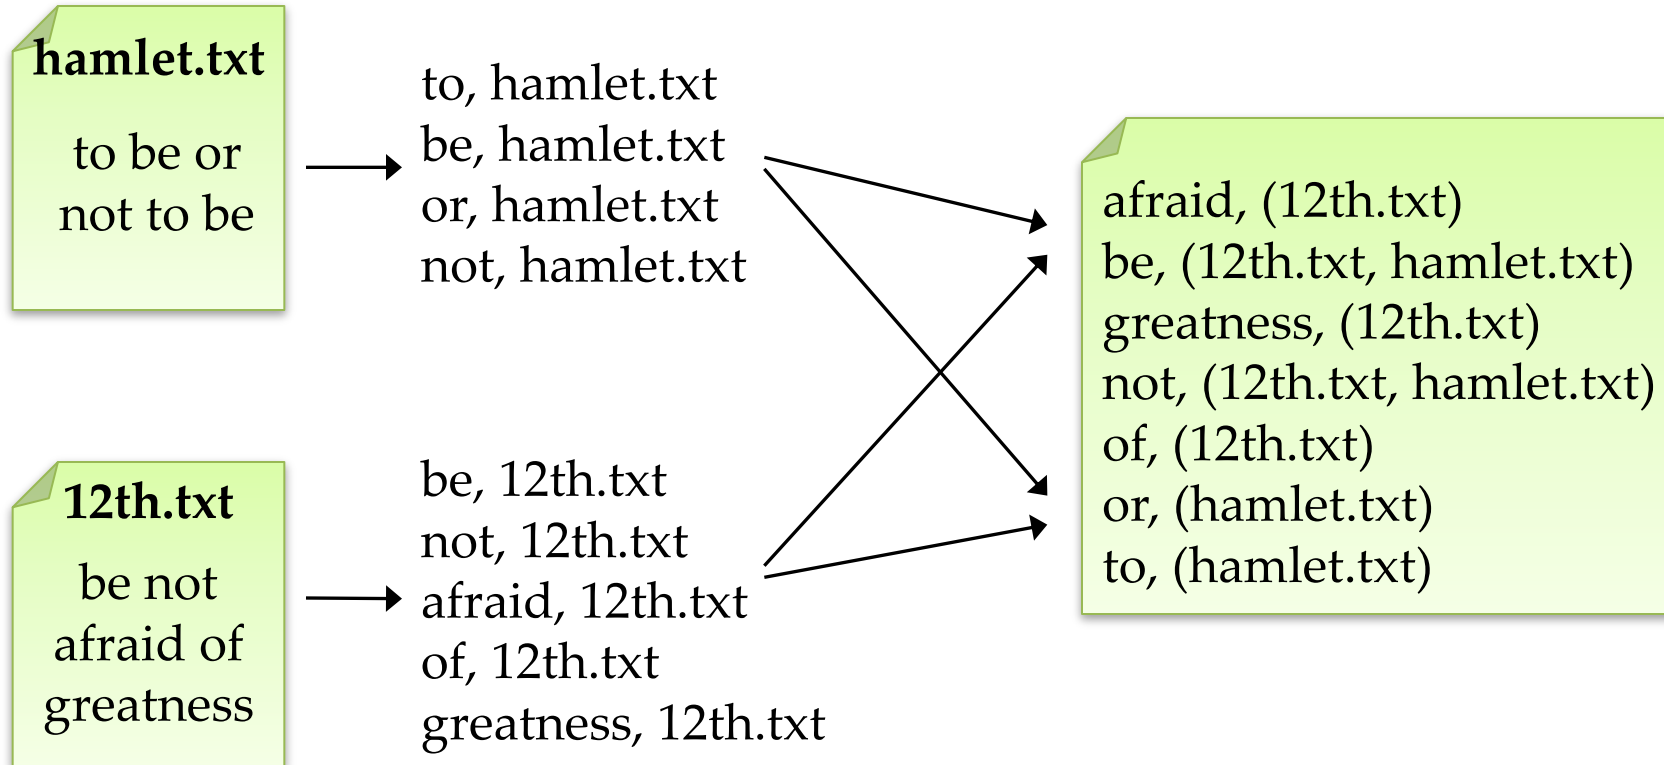
```
foreach word in text.split():  
    output(word, filename)
```

- **Combine:** uniquify filenames for each word

- **Reduce:**

```
def reduce(word, filenames):  
    output(word, sort(filenames))
```

Inverted Index Example



4. Most Popular Words

- **Input:** (filename, text) records
- **Output:** the 100 words occurring in most files
- Two-stage solution:
 - **Job 1:**
 - Create inverted index, giving (word, list(file)) records
 - **Job 2:**
 - Map each (word, list(file)) to (count, word)
 - Sort these records by count as in sort job

5. Numerical Integration

- **Input:** (start, end) records for sub-ranges to integrate
 - Can implement using custom InputFormat
- **Output:** integral of $f(x)$ over entire range

- **Map:**

```
def map(start, end):  
    sum = 0  
    for(x = start; x < end; x += step):  
        sum += f(x) * step  
    output("", sum)
```

- **Reduce:**

```
def reduce(key, values):  
    output(key, sum(values))
```

Hints for Map Reduce Programming

- Decide the way to partition the original data
- Avoid skewness
- Try to reduce the workload of network transmission